

A Nonasymptotic Theory of Gain-Dependent Error Dynamics in Behavior Cloning

Junghoon Seo

Abstract—Behavior cloning (BC) policies on position-controlled robots inherit the closed-loop response of the underlying PD controller, yet the nonasymptotic finite-horizon consequences of controller gains for BC failure remain open. We show that independent sub-Gaussian action errors propagate through the gain-dependent closed-loop dynamics to yield sub-Gaussian position errors whose proxy matrix $X_\infty(K)$ governs the failure tail. The probability of horizon- T task failure factorizes into a gain-dependent amplification index $\Gamma_T(K)$ and the validation loss plus a generalization slack, so training loss alone cannot predict closed-loop performance. Under shape-preserving *upper-bound* structural assumptions the proxy admits the scalar bound $X_\infty(K) \preceq \Psi(K)\bar{X}$ with $\Psi(K)$ decomposed into label difficulty, injection strength, and contraction, ranking the four canonical regimes with compliant-overdamped (CO) tightest, stiff-underdamped (SU) loosest, and the stiff-overdamped versus compliant-underdamped ordering system-dependent. For the canonical scalar second-order PD system the closed-form continuous-time stationary variance $X_\infty^c(\alpha, \beta) = \sigma^2\alpha/(2\beta)$ is strictly monotone in stiffness and damping over the *entire* stable orthant, covering both underdamped and overdamped regimes, and the exact zero-order-hold (ZOH) discretization inherits this monotonicity. The analysis gives a nonasymptotic finite-horizon extension of the gain-dependent error-attenuation explanation in [1].

Index Terms. Behavior cloning, controller gains, sub-Gaussian concentration, imitation learning.

I. INTRODUCTION

Position-controlled action spaces have become the standard interface for learned manipulation policies [2, 3, 4]. At each step a policy $\hat{\pi}_K$ outputs a position target $a_t = \hat{\pi}_K(s_t) \in \mathbb{R}^n$ that a PD controller [5] converts to motor torques via

$$\tau_t = K_p(a_t - q_t) - K_d\dot{q}_t + g(q_t), \quad (1)$$

where $q_t, \dot{q}_t \in \mathbb{R}^n$ are the measured joint positions and velocities, $K_p, K_d \in \mathbb{R}^{n \times n}$ are the diagonal positive-definite proportional and derivative gain matrices, and $g(q_t) \in \mathbb{R}^n$ is gravity compensation. Despite the ubiquity of this interface, the choice of gains (K_p, K_d) for *learning* pipelines remains largely heuristic.

Recent work by Bronars et al. [1] systematically tunes controller gains across BC, reinforcement learning, and sim-to-real transfer. The salient finding is that BC benefits from *compliant* (low K_p) and *overdamped* (high K_d) gains. Policies trained under these settings achieve higher closed-loop success rates despite exhibiting *higher* validation loss, a contrast that Bronars et al. explain through the error-attenuation properties of the closed-loop PD controller and a

derived stationary variance. Their open-loop noise-injection experiments isolate this mechanism empirically. What remains open is a nonasymptotic finite-horizon theory that connects the same mechanism to validation loss, generalization slack, and failure probabilities. Parallel theoretical work identifies fundamental challenges in continuous-action imitation learning, including sensitivity to function approximation error [6] and noise-driven amplification during autoregressive rollout [7], but these lines do not model the controller through which actions are physically executed.

The classical compounding-error analysis pioneered by Ross and Bagnell [8] and extended in DAgger [9] bounds the policy-versus-expert performance gap by the horizon and the per-step error, treating the dynamics as a black box. Spencer et al. [10] characterize three regimes of covariate shift, again without reference to the gain-dependent structure of the closed loop. In control theory, the propagation of stochastic disturbances through linear systems is well understood via Lyapunov equations and $\mathcal{H}_2/\mathcal{H}_\infty$ norms [11], and recent work on learning-based control [12, 13] connects sample complexity and regret to system-theoretic quantities. These analyses target reinforcement learning or adaptive control, however, not the offline BC pipeline where a fixed policy is deployed open-loop with respect to the expert.

We bridge these two lines of research. The key distinction from classical compounding-error analysis [8, 9] is that our bounds depend explicitly on the controller gains through the closed-loop matrices A_K, B_K , and reveal that the relevant quantity for closed-loop performance is not the prediction loss alone but its product with a gain-dependent amplification.

Our contributions are fourfold as follows:

- 1) **Sub-Gaussian propagation** (Theorem 1). Independent sub-Gaussian action errors yield a sub-Gaussian position error whose proxy matrix $X_t(K)$ is the gain-dependent Lyapunov sum.
- 2) **Failure bound via validation loss** (Theorem 3). The probability of horizon- T task failure factorizes into a gain-dependent amplification index $\Gamma_T(K)$ and the validation loss plus generalization slack, showing that higher validation loss can coexist with a lower failure bound.
- 3) **Scalar ordering and regime comparison** (Theorems 5 and 7). Under shape-preserving *upper-bound* assumptions, the proxy obeys $X_\infty(K) \preceq \Psi(K)\bar{X}$ with $\Psi(K) = b(K)l(K)/(1 - \rho_*(K)^2)$, and monotonicity of Ψ ranks the four canonical regimes.
- 4) **Global canonical monotonicity** (Theorem 9). For the

canonical scalar second-order PD system, $X_\infty^\circ(\alpha, \beta) = \sigma^2 \alpha / (2\beta)$ is strictly monotone in stiffness and damping over the *entire* stable orthant, and the exact ZOH discretization inherits this monotonicity.

II. RELATED WORKS

a) Behavior cloning theory: Behavior cloning [14, 15] learns a policy by supervised regression on expert state-action pairs. The foundational analysis of Ross and Bagnell [8] shows that the expected cost of the learned policy grows quadratically with horizon T from compounding errors, motivating interactive approaches such as DAgger [9]. Subsequent work refines these bounds under structural assumptions [10, 16], sharpens the horizon dependence [17], and links low-level dynamical stability to task-level guarantees [18]. Building on the controller-level mechanism identified in [1], our analysis treats the controller gain as a structural variable in nonasymptotic BC error propagation.

b) Impedance control: The interaction between compliance and task performance has a long history in robotics [19, 20, 21], with variable impedance learning [22, 23, 24] adapting compliance during execution. Our contribution is orthogonal. We analyze how *fixed* gains modulate the propagation of prediction errors during BC rollout.

c) Controller design for robot learning: The choice of action space and controller gains has long been recognized as a critical design axis for robot learning [25, 26, 27]. Kim et al. [28] show that torque-level control improves task-agnostic transfer, and Bronars et al. [1] provide a systematic study of how PD gains affect BC, RL, and sim-to-real transfer, identifying compliant-overdamped settings as the BC optimum. They also propose and test a mechanistic explanation based on closed-loop error attenuation, using open-loop noise-injection experiments to isolate the effect and deriving the stationary variance that captures it. The theoretical analysis below extends that account to sub-Gaussian propagation, finite-horizon failure bounds, and generalization slack.

III. PROBLEM FORMULATION

A. Gain-Dependent Closed-Loop Error Dynamics

Consider a robot with n joints and inertia $M \in \mathbb{R}^{n \times n}$ operating under the PD controller (1) with gain setting $K = (K_p, K_d)$. Let q_t^* denote the expert joint trajectory and a_t^* the expert action. Define the position error $e_t := q_t - q_t^*$ and the prediction error $\xi_t := \hat{\pi}_K(s_t) - a_t^*$. Substituting $a_t = a_t^* + \xi_t$ into (1) and linearizing about the expert trajectory yields

$$M \ddot{e}_t = -K_p e_t - K_d \dot{e}_t + K_p \xi_t, \quad (2)$$

so the action error enters scaled by K_p and stiffer gains amplify both the restoring force on e_t and the injection of action errors. Setting $x_t := [e_t^\top, \dot{e}_t^\top]^\top \in \mathbb{R}^{2n}$ gives the continuous state-space form

$$\dot{x}_t = A_K^c x_t + B_K^c \xi_t, \quad e_t = C x_t, \quad (3)$$

with $A_K^c = \begin{bmatrix} 0 & I_n \\ -M^{-1}K_p & -M^{-1}K_d \end{bmatrix}$, $B_K^c = \begin{bmatrix} 0 \\ M^{-1}K_p \end{bmatrix}$, and $C = [I_n, 0]$.

The policy emits a new target every Δt seconds and the controller holds it constant between updates, so the sampled state obeys

$$x_{t+1} = A_K x_t + B_K \xi_t, \quad e_t = C x_t, \quad (4)$$

with matrices given by the *exact zero-order-hold (ZOH)* sampling of (3),

$$A_K = e^{A_K^c \Delta t}, \quad B_K = \left(\int_0^{\Delta t} e^{A_K^c s} ds \right) B_K^c. \quad (5)$$

We consider gains for which the sampled closed loop is asymptotically stable, $\rho(A_K) < 1$.

B. Action Prediction Error Model

We model the prediction errors $\{\xi_t\}$ as *independent, mean-zero, sub-Gaussian* random vectors. There exists a PSD matrix $\Sigma_K^{\text{roll}} \succeq 0$ such that

$$\mathbb{E}[\exp(\lambda^\top \xi_t)] \leq \exp\left(\frac{1}{2} \lambda^\top \Sigma_K^{\text{roll}} \lambda\right), \quad \forall \lambda \in \mathbb{R}^m. \quad (6)$$

The proxy covariance Σ_K^{roll} governs the tail behavior, with the subscript K reflecting that different gains induce different action labels and hence different residual statistics. Independence is justified in the early-rollout regime before significant covariate shift, and sub-Gaussianity accommodates bounded errors and light-tailed neural network residuals. The *rollout-local population action MSE* is $L_{\text{roll}}(K) := \mathbb{E}[\|\xi_t\|_2^2] = \text{tr}(\Sigma_K^{\text{roll}})$.

C. Task Success, Validation Loss, and Amplification

A rollout is successful if the position error stays within a success tube of radius $r > 0$ over horizon T . The complementary failure event is

$$\text{Fail}_T := \{\exists t \leq T : \|e_t\|_2 \geq r\}. \quad (7)$$

After training, the empirical validation loss over N_{va} held-out pairs $\{(\tilde{s}_j, \tilde{a}_j)\}$ is $\hat{L}_{\text{va}}(K) := N_{\text{va}}^{-1} \sum_j \|\hat{\pi}_K(\tilde{s}_j) - \tilde{a}_j\|_2^2$, and the training loss $\hat{L}_{\text{tr}}(K)$ is defined analogously. A standard generalization argument gives, with probability at least $1 - \delta$,

$$L_{\text{roll}}(K) \leq \hat{L}_{\text{va}}(K) + \varepsilon_{\text{gen}}(K, \delta). \quad (8)$$

We define the *finite-horizon amplification index*

$$\Gamma_T(K) := \max_{0 \leq t \leq T} \sum_{s=0}^{t-1} \|C A_K^s B_K\|_{\text{op}}^2, \quad (9)$$

which measures the worst-case cumulative operator-norm gain from action errors to position errors and depends on the gains through A_K and B_K .

IV. MAIN RESULTS

We present four results in a logical chain: a sub-Gaussian propagation theorem, a closed-loop failure bound, a scalar ordering reduction in the upper-bound direction, and a regime comparison.

A. Sub-Gaussian Error Propagation

Theorem 1 lifts the per-step sub-Gaussian assumption on the action error to a horizon-wide statement about the position error. The lift is a direct consequence of linearity: a weighted sum of independent sub-Gaussians is sub-Gaussian, and the weights are exactly the closed-loop impulse response at lag s . The Lyapunov limit $S_\infty(K)$ is the standard quantity for disturbance-to-state propagation in discrete LTI systems, and the role of this paper is to track how every ingredient depends on the controller gain, so tuning (K_p, K_d) reshapes the entire tail of e_t .

Theorem 1 (Sub-Gaussian propagation). *Consider (4) with $x_0 = 0$ and action errors satisfying (6). Define the proxy matrix*

$$S_t(K) := \sum_{s=0}^{t-1} A_K^s B_K \Sigma_K^{\text{roll}} B_K^\top (A_K^\top)^s, \quad (10)$$

and its position-error projection $X_t(K) := C S_t(K) C^\top$. Then e_t is mean-zero sub-Gaussian with proxy $X_t(K)$. For all $u \in \mathbb{R}^n$,

$$\mathbb{E}[\exp(u^\top e_t)] \leq \exp\left(\frac{1}{2} u^\top X_t(K) u\right), \quad (11)$$

$$\mathbb{P}(|u^\top e_t| \geq r) \leq 2 \exp\left(-\frac{r^2}{2 u^\top X_t(K) u}\right). \quad (12)$$

If $\rho(A_K) < 1$, then $S_t(K) \rightarrow S_\infty(K)$, the unique PSD solution of the discrete Lyapunov equation

$$S_\infty(K) = A_K S_\infty(K) A_K^\top + B_K \Sigma_K^{\text{roll}} B_K^\top, \quad (13)$$

and the stationary proxy is $X_\infty(K) = C S_\infty(K) C^\top$.

Proof. Unrolling (4) from $x_0 = 0$ gives $e_t = \sum_{s=0}^{t-1} C A_K^s B_K \xi_{t-1-s}$. For $u \in \mathbb{R}^n$ and $h_s := B_K^\top (A_K^\top)^s C^\top u$, we have $u^\top e_t = \sum_s h_s^\top \xi_{t-1-s}$. By independence and the sub-Gaussian MGF bound,

$$\begin{aligned} \mathbb{E}[e^{u^\top e_t}] &\leq \prod_s \exp\left(\frac{1}{2} h_s^\top \Sigma_K^{\text{roll}} h_s\right) \\ &= \exp\left(\frac{1}{2} u^\top X_t(K) u\right). \end{aligned} \quad (14)$$

The tail bound follows from the Chernoff method, and Schur stability with the discrete Lyapunov theorem yields convergence to $X_\infty(K)$. $\square$

Corollary 2 (Euclidean tail bound). *Under Theorem 1,*

$$\mathbb{P}(\|e_t\|_2 \geq r) \leq 2n \exp\left(-\frac{r^2}{2n \lambda_{\max}(X_t(K))}\right). \quad (15)$$

Proof. Diagonalize $X_t = Q \Lambda Q^\top$ and set $y = Q^\top e_t$. If $\|y\|_2 \geq r$ then some $|y_i| \geq r/\sqrt{n}$, and each y_i is sub-Gaussian with proxy $\lambda_i \leq \lambda_{\max}(X_t)$. A union bound over coordinates gives the result. $\square$

The dimensional factor n enters twice, once as a linear prefactor and once inside the denominator of the exponent. In the scalar case $n = 1$ both factors disappear and the bound reduces to the one-dimensional Chernoff tail. For robot manipulators n is typically between six and twelve, and the corresponding slack is mild on a logarithmic scale.

B. Closed-Loop Failure Bound

Theorem 3 (Failure bound via validation loss). *Under the framework of Section III, with probability at least $1 - \delta$ over training randomness,*

$$\begin{aligned} \mathbb{P}(\text{Fail}_T \mid \hat{\pi}_K) &\leq \\ 2n(T+1) \exp\left(-\frac{r^2}{2n \Gamma_T(K) (\hat{L}_{\text{va}}(K) + \varepsilon_{\text{gen}}(K, \delta))}\right). \end{aligned} \quad (16)$$

Proof. From Theorem 1, $\lambda_{\max}(X_t(K)) \leq \sum_s \|C A_K^s B_K\|_{\text{op}}^2 \lambda_{\max}(\Sigma_K^{\text{roll}}) \leq \Gamma_T(K) L_{\text{roll}}(K)$, where the second inequality uses $\lambda_{\max}(\Sigma_K^{\text{roll}}) \leq \text{tr}(\Sigma_K^{\text{roll}}) = L_{\text{roll}}(K)$. The generalization bound (8) replaces $L_{\text{roll}}(K)$ by $\hat{L}_{\text{va}}(K) + \varepsilon_{\text{gen}}(K, \delta)$ with probability $1 - \delta$. Substituting into (15) and union-bounding over $t = 0, 1, \dots, T$ yields (16). $\square$

The exponent of (16) is governed by the product $\Gamma_T(K) (\hat{L}_{\text{va}}(K) + \varepsilon_{\text{gen}}(K, \delta))$ rather than by the loss alone, so a gain K_1 with lower training or validation loss than K_2 can incur a strictly worse failure bound once its amplification is correspondingly larger. This matches the attenuation-difficulty tradeoff of [1], in which compliant-overdamped gains attain higher closed-loop success despite higher validation MSE because their amplification is small enough to dominate the product. Appendix A of [1] derives the stationary variance $\sigma^2 K_p / (2K_d)$, and its *State-Space Impact of Policy Errors* corollary uses $\sqrt{K_p / (2K_d)} \epsilon_\pi(K)$ to explain the same mechanism; Theorem 3 extends that analysis to a nonasymptotic finite-horizon tail bound carrying the explicit generalization slack $\varepsilon_{\text{gen}}(K, \delta)$.

C. Scalar Ordering Reduction

To compare gain regimes through a single scalar, we impose structural conditions on the error dynamics that yield an *upper bound* on the proxy matrix. Every inequality below is an upper bound, and reversing any one invalidates the chain.

Definition 4 (Shape-preserving upper-bound structure). *The error dynamics (4) have shape-preserving upper-bound structure if there exist fixed PSD matrices $\bar{W}, \bar{\Sigma} \succeq 0$ and scalar functions $l(K), b(K) > 0$, $\rho_*(K) \in [0, 1]$ such that*

$$\begin{aligned} \text{(i) label difficulty:} & \quad \Sigma_K^{\text{roll}} \preceq l(K) \bar{\Sigma}, \\ \text{(ii) injection:} & \quad B_K \bar{\Sigma} B_K^\top \preceq b(K) \bar{W}, \\ \text{(iii) Lyapunov contraction:} & \quad A_K \bar{W} A_K^\top \preceq \rho_*(K)^2 \bar{W}. \end{aligned}$$

Condition (i) bounds the rollout-error proxy at gain K , condition (ii) bounds how strongly action errors are injected into the state space, and condition (iii) certifies $\bar{W}$ as a Lyapunov function for A_K with rate $\rho_*(K)^2$.

Theorem 5 (Scalar ordering reduction). *Under Definition 4, define the ordering index*

$$\Psi(K) := \frac{b(K) l(K)}{1 - \rho_*(K)^2}. \quad (17)$$

Then

$$X_\infty(K) \preceq \Psi(K) \bar{X}, \quad \bar{X} := C \bar{W} C^\top. \quad (18)$$

Consequently, $\lambda_{\max}(X_\infty(K)) \leq \Psi(K) \lambda_{\max}(\bar{X})$. If $\Psi(K_1) \leq \Psi(K_2)$, then K_1 's proxy upper bound is no larger than K_2 's, and substituting into (16) yields a failure-probability upper bound for K_1 that is no larger than the corresponding upper bound for K_2 .

Proof. By (i), left- and right-multiplying $\Sigma_K^{\text{roll}} \preceq l(K) \bar{\Sigma}$ by B_K and its transpose and applying (ii) gives $B_K \Sigma_K^{\text{roll}} B_K^\top \preceq l(K) b(K) \bar{W}$. The map $X \mapsto A_K^s X (A_K^\top)^s$ preserves $\preceq$ for $s \geq 0$, so $A_K^s B_K \Sigma_K^{\text{roll}} B_K^\top (A_K^\top)^s \preceq l(K) b(K) A_K^s \bar{W} (A_K^\top)^s$. Iterating (iii) yields $A_K^s \bar{W} (A_K^\top)^s \preceq \rho_*(K)^{2s} \bar{W}$, and the geometric sum gives $S_\infty(K) \preceq l(K) b(K) \sum_{s \geq 0} \rho_*(K)^{2s} \bar{W} = \Psi(K) \bar{W}$. Projecting through C yields (18), and operator monotonicity of $\lambda_{\max}$ followed by substitution into (16) gives the failure-bound monotonicity. $\square$

The index $\Psi(K)$ assembles three independent penalties: a large label difficulty $l(K)$, a large injection $b(K)$, and slow contraction (large $\rho_*(K)$). A gain can remain optimal even when its label difficulty is the highest of the four candidates, provided the injection and contraction terms are correspondingly smaller. This three-way competition is precisely what makes the BC gain-tuning problem nontrivial, and Theorem 5 reduces the resolution of that competition to a single scalar comparison. Importantly, the three penalties are not symmetric in the gains: the injection $b(K)$ is typically a strictly increasing function of K_p (stiffer gains inject action errors more aggressively), while the contraction $\rho_*(K)$ is typically a strictly decreasing function of K_d (stronger damping tightens the Lyapunov rate). The label-difficulty term $l(K)$ is the only component with a qualitatively data-driven behavior: harder labels are induced by compliant settings because the policy must compensate for richer closed-loop trajectories.

Corollary 6 (Finite-horizon convergence). *Under Definition 4, the truncation error of the proxy admits the upper bound*

$$X_\infty(K) - X_t(K) \preceq \frac{\rho_*(K)^{2t}}{1 - \rho_*(K)^2} b(K) l(K) \bar{X}, \quad (19)$$

and consequently $\|X_\infty(K) - X_t(K)\|_{\text{op}} \leq \rho_*(K)^{2t} \Psi(K) \|\bar{X}\|_{\text{op}}$, so the approximation error decays geometrically at rate $\rho_*(K)^2$ per step.

Proof. $S_\infty(K) - S_t(K) = \sum_{s=t}^{\infty} A_K^s W_K (A_K^\top)^s$ with $W_K = B_K \Sigma_K^{\text{roll}} B_K^\top$. Repeating the upper-bound chain from the proof of Theorem 5 on the tail sum gives $X_\infty(K) - X_t(K) \preceq l(K) b(K) \rho_*(K)^{2t} / (1 - \rho_*(K)^2) \bar{X}$. The operator-norm bound follows since $\|\cdot\|_{\text{op}}$ is monotone under $\preceq$ for PSD matrices. $\square$

D. Regime Ordering

We parameterize gains as $K = (\alpha, \beta)$ where α denotes stiffness (K_p scale) and β denotes damping (K_d scale), and




Fig. 1. Hasse diagram of the regime ordering established by Theorem 7. Solid arrows indicate $\Psi(K_1) \leq \Psi(K_2)$. CO is the unique minimum and SU the unique maximum. The dashed line between SO and CU indicates that their ordering is system-dependent.

define four canonical regimes

$$\begin{aligned} \text{CO} &= (\alpha_L, \beta_H), & \text{SO} &= (\alpha_H, \beta_H), \\ \text{CU} &= (\alpha_L, \beta_L), & \text{SU} &= (\alpha_H, \beta_L), \end{aligned} \quad (20)$$

with $\alpha_H > \alpha_L$ and $\beta_H > \beta_L$.

Theorem 7 (Regime ordering). *Suppose $\Psi(\alpha, \beta)$ is non-decreasing in α and non-increasing in β . Then*

$$\Psi(\text{CO}) \leq \min\{\Psi(\text{SO}), \Psi(\text{CU})\}, \quad (21)$$

$$\Psi(\text{SU}) \geq \max\{\Psi(\text{SO}), \Psi(\text{CU})\}, \quad (22)$$

so CO admits the tightest failure-bound upper bound and SU the loosest. The SO vs. CU comparison reduces to $\Psi(\alpha_H, \beta_H) \leq \Psi(\alpha_L, \beta_L)$ and is therefore system-dependent.

Proof. Direct from monotonicity. $\Psi(\alpha_L, \beta_H) \leq \Psi(\alpha_H, \beta_H)$ by stiffness-monotonicity, and $\Psi(\alpha_L, \beta_H) \leq \Psi(\alpha_L, \beta_L)$ by damping-antimonotonicity. Similarly $\Psi(\alpha_H, \beta_H) \leq \Psi(\alpha_H, \beta_L)$ and $\Psi(\alpha_L, \beta_L) \leq \Psi(\alpha_H, \beta_L)$. The SO-CU comparison involves opposing effects and cannot be resolved without additional information. $\square$

E. Verification for Canonical Second-Order Systems

Theorem 7 reduces the regime ordering to monotonicity of $\Psi(K)$. We verify this monotonicity *directly* for the canonical second-order PD system without invoking the shape-preserving abstraction. The continuous-time analysis yields a closed-form expression that is monotone over the entire stable gain space, covering both underdamped and overdamped regimes, and the exact ZOH sampling used in Section V inherits the same monotonicity.

Lemma 8 (Continuous-time stationary covariance). *For the scalar ($n = 1$) continuous-time error dynamics (3) with mass m and gains $K_p = \alpha$, $K_d = \beta$ driven by a*

white action-error process of intensity σ^2 , the stationary covariance of x_t is the unique PSD solution of $A_K^c P + P(A_K^c)^\top + B_K^c \sigma^2 (B_K^c)^\top = 0$, given by

$$P^c(\alpha, \beta) = \begin{bmatrix} \frac{\sigma^2 \alpha}{2\beta} & 0 \\ 0 & \frac{\sigma^2 \alpha^2}{2\beta m} \end{bmatrix}. \quad (23)$$

The position-error stationary variance is therefore

$$X_\infty^c(\alpha, \beta) = C P^c C^\top = \frac{\sigma^2 \alpha}{2\beta}. \quad (24)$$

Proof. Write $A_c = \begin{bmatrix} 0 & 1 \\ -\alpha/m & -\beta/m \end{bmatrix}$, $B_c = (\alpha/m)e_2$, and $P = \begin{bmatrix} p_{11} & p_{12} \\ p_{12} & p_{22} \end{bmatrix}$. Expanding $A_c P + P A_c^\top + B_c \sigma^2 B_c^\top = 0$ entry-wise gives the three independent equations $2p_{12} = 0$, $p_{22} = (\alpha/m)p_{11}$, and $-2(\beta/m)p_{22} + (\alpha/m)^2 \sigma^2 = 0$, solved in order to obtain $p_{12} = 0$, $p_{22} = \sigma^2 \alpha^2 / (2\beta m)$, and $p_{11} = \sigma^2 \alpha / (2\beta)$. Hurwitz stability of A_c gives uniqueness of the PSD solution. $\square$

Theorem 9 (Global monotonicity in canonical systems). *The scalar canonical stationary variance $X_\infty^c(\alpha, \beta) = \sigma^2 \alpha / (2\beta)$ is strictly increasing in stiffness α and strictly decreasing in damping β over the entire stable region $\{(\alpha, \beta) \in \mathbb{R}^2 : \alpha, \beta > 0\}$, covering both underdamped ($\beta^2 < 4m\alpha$) and overdamped ($\beta^2 \geq 4m\alpha$) regimes. The canonical regime ordering of Theorem 7 therefore holds with $\Psi^c(\alpha, \beta) := X_\infty^c(\alpha, \beta) / \bar{X}^c$ for any positive scalar reference $\bar{X}^c$.*

Proof. From (24), $\partial_\alpha X_\infty^c = \sigma^2 / (2\beta) > 0$ and $\partial_\beta X_\infty^c = -\sigma^2 \alpha / (2\beta^2) < 0$ on the positive orthant. The Hurwitz condition for the canonical scalar system is exactly $\alpha, \beta > 0$, so the closed form covers the entire stable region without an underdamped/overdamped split, and the four regimes inherit the ordering of Theorem 7. $\square$

Working in continuous time keeps the Lyapunov solution in closed form and avoids any branching on the underdamped/overdamped split, while inheritance through exact ZOH discretization matches the input matrix used throughout Section V. The scalar $\sigma^2 K_p / (2K_d)$ is the formal stationary-variance quantity derived in the analytical appendix of [1] under a one-degree-of-freedom Gaussian white-noise model. Lemma 8 recovers the same quantity from the continuous-time Lyapunov equation, and Theorem 9 together with Proposition 10 extends the identity to sub-Gaussian forcing and to strict monotonicity on the entire stable orthant under exact ZOH.

Proposition 10 (Discrete-time inheritance under ZOH). *For the canonical scalar system sampled by (5) with control period $\Delta t > 0$, the discrete stationary proxy $X_\infty^d(\alpha, \beta, \Delta t)$ is smooth in (α, β) on the open Schur-stable region, and $X_\infty^d(\alpha, \beta, \Delta t) / \Delta t \rightarrow X_\infty^c(\alpha, \beta)$ as $\Delta t \rightarrow 0^+$. For control periods small enough that the leading-order continuous-time term dominates, X_∞^d inherits the global strict monotonicity of X_∞^c in (α, β) , covering the gain ranges used at typical BC control rates of 50–1000 Hz.*

Proof. The ZOH matrices $A_K(\Delta t) = e^{A_K^c \Delta t}$ and $B_K(\Delta t) = (\int_0^{\Delta t} e^{A_K^c s} ds) B_K^c$ are smooth in $(\alpha, \beta, \Delta t)$ on the stable region, with $A_K = I + A_K^c \Delta t + O(\Delta t^2)$ and $B_K = \Delta t B_K^c + O(\Delta t^2)$. Substituting into (13) and expanding to leading order yields $A_K^c S_\infty^d + S_\infty^d (A_K^c)^\top + \Delta t B_K^c \sigma^2 (B_K^c)^\top + O(\Delta t^2) = 0$, so $S_\infty^d / \Delta t \rightarrow P^c$. Strict inequalities are open under continuous perturbation, so the strict monotonicity of X_∞^c from Theorem 9 carries over, and the numerical sweep of Fig. 3 confirms it on the experimental gain range at $\Delta t = 0.02$ s. $\square$

Corollary 11 (Explicit failure bound, canonical case). *For the canonical scalar second-order system, the failure-bound exponent of Theorem 3 admits the explicit form*

$$\mathbb{P}(\text{Fail}_T) \leq 2(T+1) \exp\left(-\frac{r^2 \beta}{\alpha \Delta t L_{\text{roll}}(K)}\right) + o(\Delta t), \quad (25)$$

where the remainder collects Δt^2 and higher discretization corrections via Proposition 10. The exponent therefore scales as β/α for fixed validation loss, so increasing damping or decreasing stiffness sharpens the bound.

Proof. Proposition 10 gives $X_\infty^d / \Delta t = X_\infty^c(\alpha, \beta) + O(\Delta t)$, and Lemma 8 identifies the leading term as $\sigma^2 \alpha / (2\beta)$. Substituting $X_\infty^d = \Delta t \sigma^2 \alpha / (2\beta) + o(\Delta t)$ into (15) with $n = 1$ and $\sigma^2 = L_{\text{roll}}(K)$ yields (25). $\square$

For a fixed prediction-error budget, equation (25) depends only on the ratio β/α . Doubling both gains leaves the bound unchanged, which decouples the gain-tuning question from the absolute scale of the controller.

Corollary 12 (Connection to the squared $\mathcal{H}_2$ norm). *For the canonical second-order system, the closed-form continuous-time stationary variance equals the squared $\mathcal{H}_2$ norm of the disturbance-to-position transfer function $H_K(s) = (\alpha/m) / (s^2 + (\beta/m)s + \alpha/m)$ scaled by σ^2 , and the amplification index of Theorem 3 converges in the long-horizon limit to $\lim_{T \rightarrow \infty} \Gamma_T(K) = \|\mathcal{G}_K\|_{\mathcal{H}_2}^2$. Minimizing $\Psi(K)$ recovers the classical $\mathcal{H}_2$ -optimal gain selection rule as the failure-bound minimizer.*

Proof. The transfer function $H_K(s)$ has natural frequency $\omega_n = \sqrt{\alpha/m}$ and damping ratio $\zeta = \beta / (2\sqrt{m\alpha})$. The standard $\mathcal{H}_2$ formula gives $\|H_K\|_{\mathcal{H}_2}^2 = (\alpha/m)^2 / (4\zeta\omega_n^3) = \alpha / (2\beta)$, matching Lemma 8 after the noise scaling. The asymptotic identification of Γ_T is standard for Schur-stable LTI systems [11]. $\square$

V. NUMERICAL ILLUSTRATION

We illustrate the framework on the canonical scalar second-order system, verifying sub-Gaussian propagation, the failure bound, and the regime ordering. The closed-loop matrices are obtained by exact ZOH sampling (5), matching Proposition 10.

A. Setup

Consider unit mass ($m = 1$ kg) under PD control sampled at $\Delta t = 0.02$ s (50 Hz). We evaluate the four canonical gain regimes of (20) with $\alpha_L = 50$, $\alpha_H = 100$, $\beta_L = 20$, $\beta_H =$




Fig. 2. Monte Carlo position-error envelopes ($N = 50,000$ rollouts) for the four gain regimes under exact ZOH discretization. Shaded bands show the 95th and 99th percentiles. The dashed line is the steady-state 95th-percentile bound from Theorem 1. All panels share the same vertical scale. CO yields the tightest envelopes, confirming the predicted ordering.




Fig. 3. Stationary proxy $X_\infty(K)$ (normalized by X_∞^{CO} , log scale) over the gain parameter space, computed via exact ZOH discretization at $\Delta t = 0.02$ s. Markers locate the four canonical regimes. The landscape is monotone increasing in K_p and decreasing in K_d over the entire stable region, consistent with Theorem 9 and Proposition 10.

TABLE I

SYSTEM QUANTITIES AND EMPIRICAL FAILURE RATES FOR THE SCALAR CANONICAL SYSTEM ($m = 1$ KG, $\Delta t = 0.02$ s, $T = 50$, $r = 0.3$). X_∞^c FROM (24), X_∞^d FROM THE ZOH DISCRETE LYAPUNOV.

	CO	SO	CU	SU
K_p	50	100	50	100
K_d	40	40	20	20
$\rho(A_K)$	0.974	0.948	0.943	0.819
$X_\infty^c = \sigma^2 K_p / (2K_d)$	0.625	1.250	1.250	2.500
X_∞^d (ZOH discrete)	0.012	0.025	0.025	0.050
$X_\infty^d / X_\infty^{d, \text{CO}}$	1.0	2.0	2.0	4.0
$\hat{P}(\text{Fail}_T)$	0.01	0.26	0.22	0.75

40. Action prediction errors are drawn i.i.d. from $\mathcal{N}(0, 1)$, satisfying (6) with $\Sigma_K^{\text{roll}} = 1$. For each regime we simulate $N = 50,000$ independent rollouts of $T = 50$ steps (1 s task horizon), and compute the stationary proxy $X_\infty(K)$ from the discrete Lyapunov equation (13).

B. Sub-Gaussian Error Propagation

Fig. 2 shows the position-error magnitude $\|e_t\|$ across regimes, with median, 95th-percentile, and 99th-percentile envelopes from the Monte Carlo ensemble. The dashed line




Fig. 4. Empirical failure rate $\hat{P}(\text{Fail}_T)$ (solid) versus the Theorem 3 upper bound (dashed) as a function of the success-tube radius r for the four canonical regimes. The bound dominates the Monte Carlo curve in every regime, and the regime ordering $\text{CO} \prec \text{SO} \approx \text{CU} \prec \text{SU}$ is preserved at all r .

is the steady-state 95th-percentile threshold predicted by Theorem 1 via $r_{95} = \sqrt{2X_\infty(K) \ln 40}$, derived from (12) with $n = 1$. The empirical 95th percentile stays inside the predicted bound in every regime, confirming sub-Gaussian propagation. The error magnitudes differ systematically: CO yields steady-state errors roughly half those of SO or CU and four times smaller than SU, despite identical action-error distributions across the four regimes.

C. Failure Bound and Regime Ordering

Table I reports both the closed-form continuous-time stationary variance $X_\infty^c = \sigma^2 \alpha / (2\beta)$ and the exact discrete ZOH proxy X_∞^d . Both quantities respect the predicted ordering $\text{CO} \prec \text{SO} \approx \text{CU} \prec \text{SU}$, with the near-equality of SO and CU exemplifying the system-dependent SO/CU comparison. The Monte Carlo failure rates corroborate the ordering: at $r = 0.3$, CO fails on 1% of rollouts, SO and CU on 26% and 22%, and SU on 75%. Notably CO has a *larger* spectral radius (0.974) than SU (0.819) yet achieves a fourfold smaller proxy because its injection is proportionally weaker, reflecting the three-way competition between label difficulty, injection strength, and contraction inside $\Psi(K)$.

Fig. 3 sweeps $X_\infty(K)$ over the (K_p, K_d) parameter space. The landscape is monotone increasing in stiffness and de-

creasing in damping over the *entire* sampled region, covering both underdamped and overdamped regimes, with CO at the global minimum and SU at the maximum, consistent with Theorem 9.

D. Test of the Failure Bound

Fig. 4 compares the empirical failure rate $\hat{P}(\text{Fail}_T)$ with the upper bound of Theorem 3 as the success-tube radius r varies. For each regime, the dashed curve plots $\min\{1, 2(T+1) \exp(-r^2/(2\Gamma_T(K)L_{\text{roll}}(K)))\}$ from Corollary 2 with $L_{\text{roll}} = 1$ and the empirical $\Gamma_T(K)$ from (9). The bound dominates the Monte Carlo curve uniformly across regimes and tube radii, with CO retaining the largest gap (well-conditioned tail) and SU sitting closest to its bound, and the relative ordering matches Theorem 7 at every r . The failure-bound exponent is therefore qualitatively correct as a comparator and quantitatively conservative for absolute prediction. The sub-Gaussian assumption is the dominant source of looseness, since L_{roll} replaces the per-coordinate proxy by its trace and the union over $T+1$ steps adds an unconditional log-factor.

VI. DISCUSSION

A. Gain-Selection Criterion

Theorem 3 turns the attenuation-difficulty tradeoff of [1] into a finite-horizon decision rule. Rather than minimizing the training or validation loss, practitioners should minimize the product $\Gamma_T(K) \cdot (\hat{L}_{\text{va}}(K) + \varepsilon_{\text{gen}}(K, \delta))$, equivalently $\Psi(K)$ when shape-preserving structure holds or $X_\infty^c(K)$ for canonical second-order systems. Theorems 7 and 9 recommend the compliant-overdamped regime in agreement with [1]. In our bound, the higher-validation-MSE yet better-closed-loop-performance pattern appears inside the product: CO incurs larger label difficulty $l(K)$ but a proportionally smaller amplification, and the exponent of (16) rewards that trade. A practical workflow follows directly. Start from compliant-overdamped gains, monitor $\Gamma_T(K) \cdot \hat{L}_{\text{va}}(K)$ rather than the training loss alone, and resolve the SO versus CU choice empirically because the ordering is fundamentally system-dependent.

B. Relation to Classical Analysis

Classical compounding-error analysis [8] bounds the performance gap by $O(\varepsilon T^2)$. Our amplification index refines the T^2 factor. For Schur-stable systems, $\Gamma_T(K) = O(1/(1 - \rho(A_K))^2)$ remains bounded as $T \rightarrow \infty$ and converges to the squared $\mathcal{H}_2$ norm $\lim_{T \rightarrow \infty} \Gamma_T(K) = \|\mathcal{G}_K\|_{\mathcal{H}_2}^2$ [11], connecting the analysis to classical robust control. Block et al. [7] quantify how training noise compounds during autoregressive rollout, while our amplification index complements that line by exposing how the controller gains modulate the same pathway through A_K and B_K . The indeterminacy of the SO versus CU comparison matches the system-dependent observations reported in [1].

C. Extensions and Limitations

The linearization about the expert trajectory is the main simplification. Nonlinear, contact-rich settings are a natural next step via contraction theory [12]. The independence assumption on ξ_t weakens once covariate shift accumulates, and in that regime the bounds remain valid after replacing Σ_K^{roll} by an upper bound on the conditional sub-Gaussian proxy. Multi-step action chunks [2, 4] enter the framework unchanged by treating a length- H chunk as a single Hm -dimensional prediction error and evaluating $\Gamma_T(K)$ at the chunk-execution rate. For multi-joint systems with non-uniform gains the matrix-level comparison through $X_\infty(K)$ replaces the scalar reduction whenever the shape-preserving structure of Definition 4 fails to hold uniformly across joints.

VII. CONCLUSION

We have presented a theoretical framework that quantifies how controller gains shape behavior cloning failure through four results: sub-Gaussian error propagation, a closed-loop failure bound that factorizes into a gain-dependent amplification index and the validation loss, a scalar ordering reduction via shape-preserving upper-bound structure, and a global monotonicity theorem for canonical second-order systems covering both underdamped and overdamped regimes. The closed-form continuous-time stationary variance $X_\infty^c = \sigma^2 K_p / (2K_d)$ recovers the stationary-variance result of [1] and extends the associated attenuation mechanism to nonasymptotic finite-horizon failure bounds. Natural extensions include nonlinear contact dynamics, co-optimization of gains with policy architecture, and tightening the discrete-time inheritance bound beyond the leading-order continuous-time limit.

APPENDIX I

ASYMPTOTIC FORM OF Γ_T FOR THE CANONICAL SYSTEM

For the canonical scalar second-order system under exact ZOH discretization with sampling period Δt , the closed-loop matrix A_K has eigenvalues

$$\lambda_{1,2}(K) = \exp\left(\frac{-\beta \pm \sqrt{\beta^2 - 4m\alpha}}{2m} \Delta t\right), \quad (26)$$

which lie strictly inside the unit disk on the open Hurwitz region $\alpha, \beta > 0$. Bounding the geometric series of operator norms in (9) yields

$$\Gamma_T(K) \leq \frac{\|B_K\|_{\text{op}}^2}{1 - \rho(A_K)^2} (1 - \rho(A_K)^{2T}), \quad (27)$$

where $\|B_K\|_{\text{op}}^2 = O(\Delta t^2)$ and $\rho(A_K)^2 = 1 - \beta \Delta t / m + O(\Delta t^2)$ as $\Delta t \rightarrow 0^+$. Substituting these expansions into (27) and using $1 - \rho^{2T} \rightarrow 1$ in the long-horizon limit gives $\Gamma_T(K) \sim \alpha \Delta t / (2\beta)$, which reproduces the leading-order failure exponent of Corollary 11 and identifies X_∞^c as the rate constant of the long-horizon failure tail.

APPENDIX II
MULTI-JOINT REDUCTION VIA DEFINITION 4

For an n -joint robot with diagonal mass $M = \text{diag}(m_i)$ and diagonal gains $K_p = \text{diag}(\alpha_i)$, $K_d = \text{diag}(\beta_i)$, the closed-loop dynamics decouple per joint. Choosing $\bar{W} := I_{2n}$ and $\bar{\Sigma} := I_n$, the conditions of Definition 4 hold with

$$\begin{aligned} l(K) &= \lambda_{\max}(\Sigma_K^{\text{roll}}), \\ b(K) &= \max_i \frac{\alpha_i^2 \Delta t^2}{4m_i^2}, \\ \rho_*(K) &= \max_i e^{-\beta_i \Delta t / (2m_i)}. \end{aligned}$$

The scalar reduction of Theorem 5 therefore ranks multi-joint gains via the joint with the worst (largest) per-joint Ψ , recovering the canonical scalar comparison in the joint-decoupled case and providing a conservative ranking when cross-joint coupling is present in Σ_K^{roll} .

REFERENCES

- [1] A. Bronars, Y. Park, and P. Agrawal, “Tune to learn: How controller gains shape robot policy learning,” *arXiv preprint arXiv:2604.02523*, 2026.
- [2] C. Chi, Z. Xu, S. Feng, E. Cousineau, Y. Du, B. Burchfiel, R. Tedrake, and S. Song, “Diffusion policy: Visuomotor policy learning via action diffusion,” *The International Journal of Robotics Research*, vol. 44, no. 10-11, pp. 1684–1704, 2025.
- [3] C. Chi, Z. Xu, C. Pan, E. Cousineau, B. Burchfiel, S. Feng, R. Tedrake, and S. Song, “Universal manipulation interface: In-the-wild robot teaching without in-the-wild robots,” *Robotics: Science and Systems*, 2024.
- [4] T. Z. Zhao, V. Kumar, S. Levine, and C. Finn, “Learning fine-grained bimanual manipulation with low-cost hardware,” *Robotics: Science and Systems*, 2023.
- [5] R. Kelly, “Pd control with desired gravity compensation of robotic manipulators: a review,” *The International Journal of Robotics Research*, vol. 16, no. 5, pp. 660–672, 1997.
- [6] M. Simchowitz, D. Pfrommer, and A. Jadbabaie, “The pitfalls of imitation learning when actions are continuous,” in *Proceedings of Thirty Eighth Conference on Learning Theory*, 2025.
- [7] A. Block, D. J. Foster, A. Krishnamurthy, M. Simchowitz, and C. Zhang, “Butterfly effects of SGD noise: Error amplification in behavior cloning and autoregression,” in *The International Conference on Learning Representations*, 2024.
- [8] S. Ross and J. A. Bagnell, “Efficient reductions for imitation learning,” *Proceedings of the Thirteenth International Conference on Artificial Intelligence and Statistics*, pp. 661–668, 2010.
- [9] S. Ross, G. J. Gordon, and J. A. Bagnell, “A reduction of imitation learning and structured prediction to no-regret online learning,” in *Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics*, 2011, pp. 627–635.
- [10] J. Spencer, S. Choudhury, A. Venkatraman, B. Ziebart, and J. A. Bagnell, “Feedback in imitation learning: The three regimes of covariate shift,” *arXiv preprint arXiv:2102.02872*, 2021.
- [11] B. Recht, “A tour of reinforcement learning: The view from continuous control,” *Annual Review of Control, Robotics, and Autonomous Systems*, vol. 2, pp. 253–279, 2019.
- [12] N. M. Boffi, S. Tu, and J.-J. E. Slotine, “Regret bounds for adaptive nonlinear control,” in *Learning for Dynamics and Control*. PMLR, 2021, pp. 471–483.
- [13] S. Tu and B. Recht, “The gap between model-based and model-free methods on the linear quadratic regulator: An asymptotic viewpoint,” in *Conference on Learning Theory*, 2019, pp. 3036–3083.
- [14] D. A. Pomerleau, “ALVINN: An autonomous land vehicle in a neural network,” *Advances in Neural Information Processing Systems*, vol. 1, 1989.
- [15] M. Bain and C. Sammut, “A framework for behavioural cloning,” in *Machine Intelligence*, 1999.
- [16] P. Florence, C. Lynch, A. Zeng, O. A. Ramirez, A. Wahid, L. Downs, A. Wong, J. Lee, I. Mordatch, and J. Tompson, “Implicit behavioral cloning,” in *Conference on Robot Learning*, 2022, pp. 158–168.
- [17] D. J. Foster, A. Block, and D. Misra, “Is behavior cloning all you need? understanding horizon in imitation learning,” *Advances in Neural Information Processing Systems*, vol. 37, 2024.
- [18] A. Block, A. Jadbabaie, D. Pfrommer, M. Simchowitz, and R. Tedrake, “Provable guarantees for generative behavior cloning: Bridging low-level stability and high-level behavior,” *Advances in Neural Information Processing Systems*, vol. 36, pp. 48 534–48 547, 2023.
- [19] N. Hogan, “Impedance control: An approach to manipulation,” in *1984 American control conference*. IEEE, 1984, pp. 304–313.
- [20] —, “Impedance control: An approach to manipulation, part ii implementation,” *Journal of dynamic systems, measurement, and control*, vol. 107, no. 1, pp. 8–16, 1985.
- [21] O. Khatib, “A unified approach for motion and force control of robot manipulators: The operational space formulation,” *IEEE Journal on Robotics and Automation*, vol. 3, no. 1, pp. 43–53, 2003.
- [22] M. Bogdanovic, M. Khadiv, and L. Righetti, “Learning variable impedance control for contact sensitive tasks,” *IEEE Robotics and Automation Letters*, 2020.
- [23] Y. Wu, F. Zhao, T. Tao, and A. Ajoudani, “A framework for autonomous impedance regulation of robots based on imitation learning and optimal control,” *IEEE Robotics and Automation Letters*, vol. 6, no. 1, pp. 127–134, 2021.
- [24] K. Kronander and A. Billard, “Learning compliant manipulation through kinesthetic and tactile human-robot interaction,” *IEEE transactions on haptics*, vol. 7, no. 3, pp. 367–380, 2013.
- [25] E. Aljalbout, F. Frank, M. Karl, and P. van der Smagt,

- “On the role of the action space in robot manipulation learning and sim-to-real transfer,” *IEEE Robotics and Automation Letters*, vol. 9, no. 6, pp. 5895–5902, 2024.
- [26] J. Eßer, G. B. Margolis, O. Urbann, S. Kerner, and P. Agrawal, “Action space design in reinforcement learning for robot motor skills,” in *Conference on Robot Learning*, 2024.
- [27] A. Mandlekar, D. Xu, J. Wong, S. Nasiriany, C. Wang, R. Kulkarni, L. Fei-Fei, S. Savarese, Y. Zhu, and R. Martín-Martín, “What matters in learning from offline human demonstrations for robot manipulation,” in *Conference on Robot Learning*, 2021.
- [28] D. Kim, G. Berseth, M. Schwartz, and J. Park, “Torque-based deep reinforcement learning for task-and-robot agnostic learning on bipedal robots using sim-to-real transfer,” *IEEE Robotics and Automation Letters*, vol. 8, no. 10, pp. 6251–6258, 2023.